

Hybrid Parallelised Framework for the solution of PDEs on HPC Clusters

By:

NIKHIL GUPTA (B16023)

RITWIK SAHA (B16110)

Mentored By:

Dr. GAURAV BHUTANI

SCEE

IIT MANDI

AIM :

To develop a hybrid framework to utilize HPC resources to solve LINEAR EQUATIONS (and hence Partial Differential Equations) using background parallelisation tools (PETSc).

Solving LINEAR EQUATIONS

$$\begin{bmatrix} \ddots & & & & \\ & -1 & 2 & -1 & \\ & & -1 & 2 & -1 \\ & & & \ddots & \end{bmatrix} \begin{bmatrix} \cdot \\ \phi^{i-1} \\ \phi^i \\ \phi^{i+1} \\ \cdot \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

ITERATIVE METHODS

- Jacobi Method
- Gauss-Seidel Method
- Steepest Descent Method
- Conjugate Gradient Method
- Many more....

PROGRESS after MID TERM EVALUATION

- Finalising the problem statement.
- Running PETSC codes on cluster.
- Benchmarking GMRES method.
- Analysing Hybrid approach for Conjugate gradient method

CONTAINERIZATION

- SINGULARITY

- To create custom runtime environment to run simulations on IIT Mandi HPC Cluster.
- Created Singularity images for both CPU and GPU implementations.
- Singularity also exposes the host's working directory inside the container unlike docker.

```

Bootstrap: docker
From: nvidia/cuda:9.0-devel

%files
  petsc-3.10.5.tar.gz

%post
  fileecho="## nvcc command
# - PATH includes /usr/local/cuda-9.0/bin
export PATH=/usr/local/cuda-9.0/bin:$PATH
# - LD_LIBRARY_PATH includes /usr/local/cuda-9.0/lib64
export LD_LIBRARY_PATH=/usr/local/cuda-9.0/lib64:$LD_LIBRARY_PATH

## Binding against the driver on the peregrine
# - LD_LIBRARY_PATH includes /usr/lib64/nvidia
export LD_LIBRARY_PATH=/usr/lib64/nvidia:$LD_LIBRARY_PATH"

  mkdir -p /usr/lib64/nvidia
  echo $fileecho > /environment
  export PATH=/usr/local/cuda-9.0/bin:$PATH
  export LD_LIBRARY_PATH=/usr/local/cuda-9.0/lib64:$LD_LIBRARY_PATH
  export LD_LIBRARY_PATH=/usr/lib64/nvidia:$LD_LIBRARY_PATH
  nvcc --version

  sed -i 's|http://archive.ubuntu|http://jp.archive.ubuntu|g' /etc/apt/sources.list
  apt update -y
  apt -y install git wget python g++ gcc gfortran curl
  apt -y install mpich libblas-dev liblapack-dev
  apt -y install build-essential valgrind
  apt -y install python3

  curl https://bootstrap.pypa.io/get-pip.py -o get-pip.py || echo "curl failed"
  python3 get-pip.py || echo "script failed"
  rm get-pip.py || echo "removal failed"
  python3 -m pip install --upgrade pip || echo "pip upgrade failed"
  python3 -m pip install matplotlib || echo "matplotlib failed"

  # wget http://ftp.mcs.anl.gov/pub/petsc/release-snapshots/petsc-3.12.1.tar.gz
  mkdir /petsc_dir/
  tar xvzf /petsc-3.10.5.tar.gz -C /petsc_dir/
  rm -r /petsc-3.10.5.tar.gz

  petsc_dir=/petsc_dir/petsc-3.10.5/


  cd $petsc_dir
  sed -i 's/^/NVCCFLAGS += -Xcompiler -openmp\nNVCCFLAGS += -Xcompiler -fopenmp\n/'
makefile
  ./configure --with-cuda=1 --with-precision=single --with-clanguage=c --download-cusp
  make all test
  cp -r ./lib/petsc /lib
  cd /

```


BENCHMARKING - I

- PETSC code
 - Solving Linear Equation for Tridiagonal Matrix
 - Using GMRES method and JACOBI preconditioner.
 - Multicore Simulation (#cores from 1 to 12)
 - GPU Simulation

OBSERVATIONS



OBSERVATIONS



CONJUGATE GRADIENT METHOD

- Our own implementation
 - Using PETSC data structures
 - Easy to choose between CPU or GPU implementation for same code
- Ran the code with **1 core CPU** implementation and **GPU**
- Implemented a **Hybrid version** of the same

CONJUGATE GRADIENT METHOD

```
VecSetRandom(x, NULL);  
MatMult(A, x, r);  
VecXPY(r, minusone, b);
```

```
//Initialize p=-r  
VecSet(p, zero);  
VecXPY(p, minusone, r);
```

```
for (i=0; i<iters; ++i){  
    VecNorm( r , NORM_2 , &rdot);  
    rdot = rdot*rdot;  
    MatMult(A, p, Ap);  
    VecDot(p, Ap, &pAp);  
  
    alpha = rdot/pAp;  
  
    VecXPY(x, alpha, p);  
  
    VecXPY(r, alpha, Ap);  
  
    VecNorm(r, NORM_2, &val);  
    val = val*val;  
    beta = val/rdot;  
  
    VecXPBY(p, minusone, beta, r);  
}
```

Compute $r_0 = Ax_0 - b$, $p_0 = -r_0$

For $k = 0, 1, 2, \dots$ **until convergence**

$$\alpha_k = \frac{r_k^T r_k}{p_k^T A p_k}$$

$$x_{k+1} = x_k + \alpha_k p_k$$


$$r_{k+1} = r_k + \alpha_k A p_k$$

$$\beta_k = \frac{r_{k+1}^T r_{k+1}}{r_k^T r_k}$$

$$p_{k+1} = -r_{k+1} + \beta_k p_k$$

End

HYBRID IMPLEMENTATION



HYBRID IMPLEMENTATION

- CPU routines for **VEC1** (n_1)
- GPU routines for **VEC2** (n_2)
- Taking maximum time of two

```
clock_t VecAXPBY_split(Vec y1, Vec y2, PetscReal alpha, PetscReal beta, Vec x1, Vec x2){  
    clock_t begin, end, a,b,r;  
  
    begin = clock();  
    VecAXPBY(y1, alpha, beta, x1);  
    end = clock();  
    a = (end - begin);  
  
    begin = clock();  
    VecAXPBY(y2, alpha, beta, x2);  
    b = (end - begin);  
  
    if( a > b ) r = a ; else r = b;  
  
    return r;  
}
```

HYBRID IMPLEMENTATION (Multi -Threaded)


- Run **CPU** and **GPU** parts on different threads, thus almost simultaneously

```
typedef struct{
    Vec x;
    Vec y;
    PetscScalar *val;
}VecDot_struct;


void* VecDot_thread(void* ptr){
    VecDot_struct *sptr = (VecDot_struct*)ptr;
    VecDot(sptr->x, sptr->y, sptr->val);
    return NULL;
}

void VecDot_split(Vec x1, Vec x2, Vec y1, Vec y2, PetscReal *value){
    PetscReal a, b;
    pthread_t t1, t2;
    VecDot_struct s1 = {x1, y1, &a}, s2 = {x2, y2, &b};
    pthread_create(&t1, NULL, VecDot_thread, (void*)&s1);
    pthread_create(&t2, NULL, VecDot_thread, (void*)&s2);
    pthread_join(t1, NULL); pthread_join(t2, NULL);
    *value = (a+b);
    return;
}
```

BENCHMARKING I - (n/2)




BENCHMARKING I - (n/2)




BENCHMARKING II - (n1, n2)

Benchmarking Hybrid Implementation




BENCHMARKING II - (n1, n2)

Benchmarking Hybrid Implementation



BENCHMARKING II - (n1, n2)



CONCLUSIONS

- PETSC is difficult to hybridise (lack of documentation)
- Multicore + GPU implementation leads to deadlock
 - Need to optimise
- Our Implementation
 - **Time(GPU) < Time (Hybrid code) < Time (CPU Code)**
- Our Aim is to achieve
 - **Time(Hybrid code) < Time(GPU)**
- IIT Mandi HPC Cluster incompatible with “Multinode + GPU” PETSC implementation.
 - Absence of InfiniBand connectivity

FUTURE SCOPE

- Improvise PETSC GPU implementation
 - Understanding PETSC structure in depth
 - New VECTOR and MATRIX type for hybridisation
- New Implementation of Linear Solvers
 - Hybrid implementation
 - CUDA / OpenMP / BLAS-LAPACK

VecType

String with the name of a PETSc vector

Synopsis

```
typedef const char* VecType;
#define VECSEQ "seq"
#define VECMPI "mpi"
#define VECSTANDARD "standard" /* seq on one process and mpi on multiple */
#define VECSHARED "shared"
#define VECSEQVIENNA "seqvienna"
#define VECMPIVIENNA "mpivienna"
#define VECVIENNA "vienna" /* seqvienna on one process and mpivienna on multiple */
#define VECSEQCUDA "seqcuda"
#define VECMPICUDA "mpicuda"
#define VECCUDA "cuda" /* seqcuda on one process and mpicuda on multiple */
#define VECNEST "nest"
#define VECNODE "node" /* use on-node shared memory */
```

<https://www.mcs.anl.gov/petsc/petsc-current/docs/manualpages/Vec/VecType.html>

THANK YOU